

Bhuvan S A

✉ bhuvansa@bhuvansa.com | ☎ [+91 8050158461](tel:+918050158461) | 🌐 bhuvansa.com | 🐱 [bhuvansa](#) | 📄 [bhuvansa](#)

Summary

AI-native full-stack engineer with 2+ years shipping production LLM systems — agentic workflows, RAG pipelines, MCP-driven tool integrations, and multi-tenant AI platforms. Built a 10-skill conversational agent on Claude Agent SDK with progressive context loading and per-user memory; redesigned an annotation pipeline to 95%+ validated accuracy. B.E. in AI/ML, daily Claude Code user, 650+ LeetCode solved.

Education

Global Academy of Technology B.E • CGPA: 8.5
Artificial Intelligence and Machine Learning Bengaluru • Dec 2021 – Jul 2025

Technical Skills

- AI / ML**
RAG, Agentic Workflows, Claude Agent SDK, MCP (Model Context Protocol), Tool/Function Calling, Structured Outputs, Prompt Engineering, Vector Search (Qdrant), Multi-Agent (crewAI), LLM Integration (Claude, OpenAI, Gemini), Whisper, Ollama
- Backend**
Python, FastAPI, Node.js, REST APIs, WebSocket, Better Auth, Flask
- Databases**
PostgreSQL, Supabase, Qdrant, Row-Level Security (RLS), Drizzle ORM, MongoDB, Redis
- Cloud & DevOps**
Docker, GitHub Actions, AWS (S3, Speech-to-Text), Azure (Blob, Container Jobs, Redis), Vercel, Git
- Frontend**
React.js, Next.js, assistant-ui, ai-sdk v6, ShadCN UI, Tailwind CSS, Radix UI
- Languages & Tools**
Python, TypeScript, JavaScript, Claude Code, Jira (Agile/Scrum)

Experience

- Compethic Remote — Oslo, Norway
Full-Stack AI Engineer Dec 2025 – Present
- Built an end-to-end multi-tenant agent system on Claude Agent SDK — 10 specialized skills with progressive context loading, Redis event bus for streaming, Azure Blob session memory, and a 140-line system prompt for skill routing and multi-skill disambiguation
 - Redesigned LLM annotation pipeline into a 3-level top-down classification tree with GPT-5 structured outputs — pruned context bloat, lowered per-review cost, and reached 95%+ categorization accuracy validated against CEO manual benchmark
 - Extended the agent to email — Microsoft Graph polling with In-Reply-To header matching, encrypted message-id → session lookup, and HTML-to-text reply parsing for cross-channel thread continuity
 - Architected the FastAPI data layer from scratch (128 endpoints, 31 routers, 80+ services) supporting the AI pipeline; parallelized annotation with asyncio + PostgreSQL FOR UPDATE SKIP LOCKED for safe concurrent multi-tenant runs

- Upsurge Labs Bengaluru
Full Stack Developer Sep 2025 – Nov 2025
- Converted REST APIs into agent-callable tool schemas for the Bhindi AI agent platform, enabling natural-language orchestration of multi-step workflows via n8n integration
 - Built a real-time chat backend from scratch — MongoDB schema, Node.js APIs, and WebSocket messaging — for the Bhindi product
 - Modernized legacy admin tooling from CommonJS to ES modules with monitoring components for platform management

Experience

Revoflo

Software Developer Intern

Remote — Bengaluru

Apr 2025 – Sep 2025

- Built a Talk-to-Database agent with MCP tools providing read-only PostgreSQL access — used transaction-scoped org context with RLS so even agent-generated SQL missing tenant filters could not leak cross-tenant data
- Integrated Qdrant vector DB for semantic search over company profiles, enabling sales reps to retrieve relevant context during calls via embedding-based similarity
- Built voice call analysis pipeline — S3 presigned upload, Whisper transcription, and AI-generated call summaries surfaced in dashboards for managers and org owners

Giftolexia Solutions

Technical Intern

Remote — Bengaluru

Mar 2023 – Jun 2024

- Built a voice transcription-based dyslexia detection system using AWS Speech-to-Text and S3, including an audio-to-text evaluation tool that compared spoken output against expected text to identify markers
- Designed a sliding window anomaly detection algorithm in Python to identify and trim stuck eye-tracking sensor segments before they reached the assessment pipeline
- Built the early screening checklist web application (Next.js, React, Flask, PostgreSQL) deployed via Docker to support clinical assessments

Projects / Open-Source

Answer Sheet Evaluator — Multi-Model AI Grading

- Architected multi-model pipeline — Gemini Flash for OCR throughput on handwritten pages and Gemini Pro for graded reasoning and personalized feedback generation
- Full-stack Next.js + PostgreSQL app with email delivery for student feedback; HOD-validated — system caught grading errors that manual review had missed

VTU Fest 2025 — Event Management Platform

- Led a team of 4 building a Next.js + PostgreSQL platform that served 4,500+ users during a college fest
- Hardened JWT auth, added SSR for performance, and code-reviewed all changes for production readiness

Honors & Awards

National Hackathon Winner — Multi-Agent Sales Bot

Feb 2024

Hack-a-League

Multi-agent system on crewAI with shared context and hand-off logic for sales conversations — agent specialization predating mainstream agentic AI

<https://www.linkedin.com/posts/activity-7160339708919582721-d6-d>

Hackathon Winner — Chat-with-PDF RAG System

Feb 2024

BrainOVision

Built complete RAG pipeline from scratch — PDF parsing, chunking, embedding generation, vector search, and context-augmented LLM responses (no LangChain)

<https://www.linkedin.com/posts/activity-7160239318727323650-P-55>

National Hackathon Winner — Dr.AIGuru

Jan 2024

Hacksprint

Privacy-first medical chatbot — locally hosted Llama 2 via Ollama with concurrent request queuing on a single laptop GPU (early local LLM serving)

https://www.linkedin.com/posts/prajwal-p-patil-5a044525a_hacksprint-hackathon-24hourhackathon-ugcPost-7150137179174768640--pW-